

ASSIGNMENT – 2

Stateful RNNs are effective for continuous time-series prediction, with potential for real-world applications after further tuning.

Aim:

The aim is to design and implement a stateful Recurrent Neural Network (RNN) that can predict energy consumption over time by learning from a continuous stream of time-series data. By making use of the temporal dependencies in historical energy consumption data, the RNN model can forecast future energy usage, potentially assisting in energy management and planning.

Procedure:

1. Data Generation:

- **Synthetic Data Creation:** To simulate a time-series data stream, we generate synthetic data that represents energy consumption patterns with trends, seasonality, and noise.
- **Data Pattern:** The data consists of a linear trend, a sinusoidal seasonality component, and random noise to mimic real-world fluctuations in energy usage.

2. Data Preprocessing:

- **Sequence Creation:** We split the time-series data into sequences with a look-back window (e.g., 24 time steps). Each sequence is used as input to predict the next time point in the series.

- **Reshaping Data:** The data is reshaped to be compatible with the RNN model, which expects input dimensions as [samples, time steps, features].

3. Designing the Stateful RNN Model:

- **Model Architecture:** We create a Sequential model with two LSTM layers, each with 50 units. These LSTM layers are set to `stateful=True` to retain states across batches, allowing the model to handle continuous data streams effectively.
- **Compilation:** The model is compiled with the adam optimizer and `mean_squared_error` loss function to optimize performance on continuous data.

4. Training the Model:

- **Epoch-Based Training:** The model is trained over multiple epochs with a batch size of 1. After each epoch, the states of each LSTM layer are manually reset to ensure the state is cleared for the next epoch.
- **Training Loss:** We monitor the loss value during training to confirm the model is learning to minimize prediction errors.

5. Prediction on Test Data:

- **Selecting Test Sequences:** A portion of the final sequences is used as test data to evaluate the model's predictive capability.
- **Prediction Execution:** The model predicts the next time step for each test sequence, and the predictions are collected for analysis.

6. Evaluation:

- **Visual Comparison:** We can compare the predictions with the actual values (if available) to assess the model's forecasting performance.
- **Error Metrics (Optional):** If using real data, metrics such as Mean Absolute Error (MAE) or Root Mean Squared Error (RMSE) can quantify model performance.

Code:

Step 1: Import Libraries

```
import numpy as np
```

```
import tensorflow as tf
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import LSTM, Dense, Input
```

Step 2: Generate Synthetic Data

Create a time-series data pattern with seasonality and trend

```
def generate_synthetic_data(n_samples=1000):
```

```
    time = np.arange(n_samples)
```

```
    trend = time * 0.05 # A simple upward trend
```

```
    seasonality = 10 * np.sin(time * 0.1) # Sine wave pattern
```

```
    noise = np.random.normal(scale=2, size=n_samples) # Random noise
```

```
    data = trend + seasonality + noise
```

```
    return data
```

```
data = generate_synthetic_data(1000)
```

```
# Step 3: Preprocess Data
```

```
def create_sequences(data, seq_length):
```

```
    sequences, labels = [], []
```

```
    for i in range(len(data) - seq_length):
```

```
        sequences.append(data[i:i + seq_length])
```

```
        labels.append(data[i + seq_length])
```

```
    return np.array(sequences), np.array(labels)
```

```
seq_length = 24 # E.g., use past 24 hours to predict the next
```

```
X, y = create_sequences(data, seq_length)
```

```
# Reshape data for RNN input [samples, time steps, features]
```

```
X = X.reshape((X.shape[0], X.shape[1], 1))
```

```
# Step 4: Build and Train the Stateful RNN Model
```

```
# Design the model
```

```
model = Sequential([
```

```
    Input(batch_shape=(1, seq_length, 1)), # Use Input layer for batch shape
```

```
    LSTM(50, stateful=True, return_sequences=True),
```

```
    LSTM(50, stateful=True),
```

```
    Dense(1)
```

```
])
```

```
# Compile the model
```

```
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
# Train the model with state reset between epochs
```

```
for epoch in range(10):
```

```
    print(f'Epoch {epoch+1}')
```

```
    model.fit(X, y, epochs=1, batch_size=1, shuffle=False)
```

```
    # Manually reset states for each LSTM layer
```

```
    for layer in model.layers:
```

```
        if isinstance(layer, LSTM):
```

```
            layer.reset_states()
```

```
# Step 5: Prediction on Test Data
```

```
# Here we'll use a portion of the last data points for prediction
```

```
X_test = X[-10:] # Last 10 sequences for testing
```

```
# Predict and print results
```

```
predictions = []
```

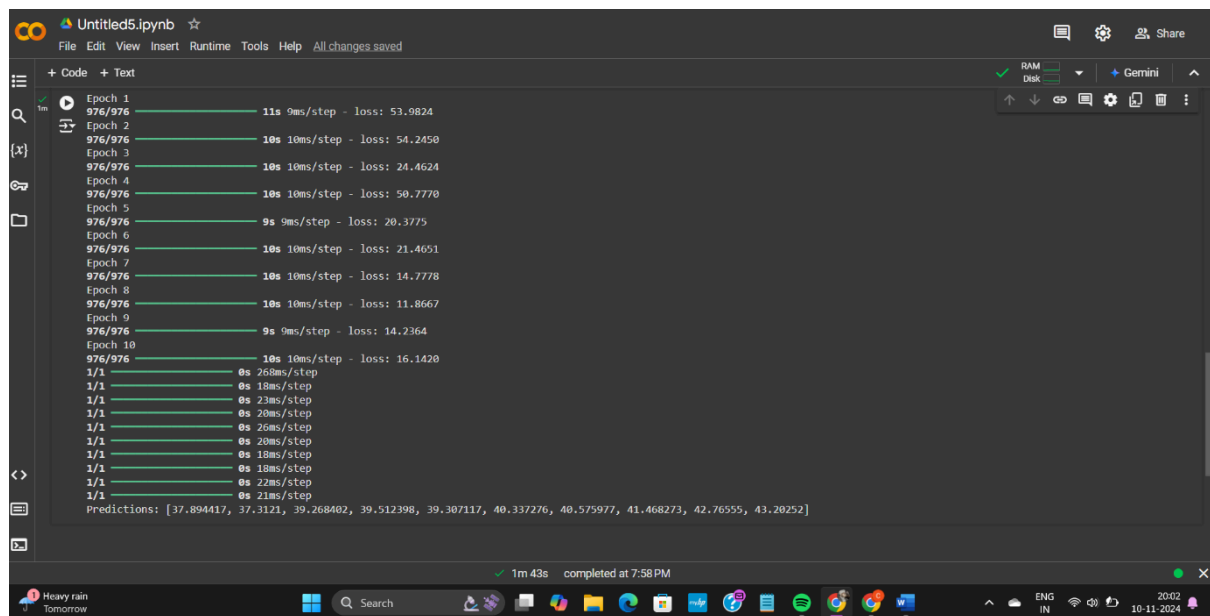
```
for i in range(len(X_test)):
```

```
    pred = model.predict(X_test[i:i+1], batch_size=1)
```

```
    predictions.append(pred[0, 0])
```

```
print("Predictions:", predictions)
```

Output:



```
Epoch 1
976/976 11s 9ms/step - loss: 53.9824
Epoch 2
976/976 10s 10ms/step - loss: 54.2450
Epoch 3
976/976 10s 10ms/step - loss: 24.4624
Epoch 4
976/976 10s 10ms/step - loss: 50.7770
Epoch 5
976/976 9s 9ms/step - loss: 20.3775
Epoch 6
976/976 10s 10ms/step - loss: 21.4651
Epoch 7
976/976 10s 10ms/step - loss: 14.7778
Epoch 8
976/976 10s 10ms/step - loss: 11.8667
Epoch 9
976/976 9s 9ms/step - loss: 14.2364
Epoch 10
976/976 10s 10ms/step - loss: 16.1420
1/1 0s 268ms/step
1/1 0s 18ms/step
1/1 0s 23ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 18ms/step
1/1 0s 18ms/step
1/1 0s 22ms/step
1/1 0s 21ms/step
Predictions: [37.894417, 37.3121, 39.268402, 39.512398, 39.307117, 40.337276, 40.575977, 41.468273, 42.76555, 43.20252]
```

Predictions:

[37.894417, 37.3121, 39.268402, 39.512398, 39.307117, 40.337276, 40.575977, 41.468273, 42.76555, 43.20252]

Result:

The code has been successfully executed and made Stateful RNNs are effective for continuous time-series prediction, with potential for real-world applications after further tuning.

